Amish Sethi

Incoming PhD Student, Harvard University

amishsethi@g.harvard.edu • amishsethi.github.io • GitHub • Google Scholar

RESEARCH INTERESTS

Reliable embodied AI, foundation models that connect video, language, and robot action, scalable neurosymbolic learning, and understanding and adapting large generative models.

EDUCATION

Harvard University — Cambridge, MA
PhD in Computer Science, incoming — Fall 2026
Advised by Heng Yang and Yilun Du. Kempner Institute Graduate Fellow and NSF Graduate Research Fellow.

University of Pennsylvania — Philadelphia, PA
B.S.E. in Computer Science and M.S.E. in Computer and Information Science — May 2026
Cumulative GPA 4.00 out of 4.00.

HONORS AND AWARDS

- NSF Graduate Research Fellowship, National Science Foundation — 2026
- Kempner Institute Graduate Fellowship, Harvard University — 2026
- John Grist Brainerd Award, Penn Engineering, top graduating senior — 2026
- CRA Outstanding Undergraduate Researcher Award, Honorable Mention — 2025
- NeurIPS 2025 Spotlight for ESCA, top 3 percent of submissions — 2025
- First Place, International Public Policy Forum — 2022
- National Merit Scholar, and ISEF Finalist in genetics research — 2021

PUBLICATIONS

Authors marked with an asterisk contributed equally. My name is shown in bold.

- **Retrieval-Augmented Vision-Language-Action Model**. Jiani Huang*, Brandon Y. Yang*, **Amish Sethi***, Yuchen Zheng, Christopher Watson, Jianing Qian, Mayur Naik, Dinesh Jayaraman. Under review at CoRL 2026. Senior thesis.
- **Do Diffusion Models Learn to Generalize Basic Visual Skills**. **Amish Sethi**, Boya Zeng, Wenhao Chai, Zhuang Liu. Under review at NeurIPS.
- **Delta Activations: A Representation for Finetuned Large Language Models**. Zhiqiu Xu*, **Amish Sethi***, Mayur Naik, Ser-Nam Lim. Under review at COLM, and published at the NeurIPS 2025 ER Workshop. arXiv.
- **ESCA: Contextualizing Embodied Agents via Scene-Graph Generation**. Jiani Huang, **Amish Sethi***, Matthew Kuo*, Mayank Keoliya, Neelay Velingker, JungHo Jung, Ziyang Li, Ser-Nam Lim, Mayur Naik. NeurIPS 2025, Spotlight, top 3 percent.
- **Dolphin: A Programmable Framework for Scalable Neurosymbolic Learning**. Aaditya Naik, Jason Liu, Claire Wang, **Amish Sethi**, Saikat Dutta, Mayur Naik, Eric Wong. ICML 2025.
- **CLAM: Unifying Finetuning, Quantization, and Pruning by Chaining LLM Adapter Modules**. Neelay Velingker, **Amish Sethi***, Jason Liu*, William Dodds*, Zhiqiu Xu, Saikat Dutta, Mayur Naik, Eric Wong. ICML 2024 ES-FoMo Workshop.
- **Functional Genetic Biomarkers of Alzheimer's Disease and Gene Expression from Peripheral Blood**. Andrew Ni*, **Amish Sethi***, Alzheimer's Disease Neuroimaging Initiative. bioRxiv 2021, cited more than 8 times.

OPEN-SOURCE AND DATASETS

- **VINE**, a promptable foundation model that turns video into spatio-temporal scene graphs, released with the ESCA-Video-87K dataset of more than 87,000 videos on Hugging Face.
- Finetuned and released more than 700 open-source language models on Hugging Face for the Delta Activations study, supporting research on reliable model ecosystems.
- The LASER codebase behind VINE reached more than 100 stars on GitHub and was featured on the Google Open Source Blog.

EXPERIENCE

AfterQuery, Research Lead Summer 2026

Leading research on the next large benchmark for embodied agents and on scaling robotics data collection well beyond teleoperation, including a more capable handheld interface for in the wild robot demonstrations.

National University of Singapore, Research Intern 2023

Built FIIGNET, a generative pipeline that synthesizes images of diseased fish for aquaponics early detection, which improved detection accuracy by 17 percent across synthetic and real datasets.

RoadBotics, Computer Vision Intern 2021

Developed a Mask R-CNN model that detects and classifies traffic signs from video with 90 percent accuracy. The system was deployed by the Pennsylvania state government to maintain an inventory of road assets.

TEACHING AND SERVICE

Head Teaching Assistant, CIS 7000 Large Language Models, Penn Fall 2024

Planned curriculum, designed assignments, held office hours, and lectured on efficient finetuning, adaptation, and evaluation for Penn's first dedicated LLM course of more than 120 students.

Undergraduate Research Mentor, Penn PURM Summer 2024

Mentored five undergraduates on the CLAM project, teaching research methodology, work with large language models, and scalable optimization.

Reviewer for the ICML 2024 ES-FoMo Workshop, AAAI 2026, and ICLR 2026. Authored a successful Grant for Faculty Mentoring Undergraduate Research, awarded 8,000 dollars by Penn for research on neurosymbolic AI.